

linear models

prashant ullegaddi
vidit bansal

July 12, 2013

linear regression

What's the problem?

- Given a set of m training examples $\{X^{(i)}, y^{(i)}\}_{i=1}^m$ where $X^{(i)} \in \mathbb{R}^{n+1}$ with $X_0^{(i)} = 1$ and $y^{(i)} \in \mathbb{R}$.
- Learn a hypothesis function $h(\cdot)$ such that for any new test example $x^{(t)}$, y can be predicted as $y^{(t)} = h(x^{(t)})$.

- Fit a hyperplane

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \cdots + \theta_n x_n = \sum_{i=0}^n \theta_i x_i = \theta^T x$$

ordinary least squares method

Optimization

Find parameters θ that:

$$\text{minimize}_{\theta} J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(X^{(i)}) - y^{(i)})^2$$

gradient descent method

Update rule

$$\theta := \theta - \alpha \nabla_{\theta} J(\theta)$$

where α is the learning rate, $\nabla_{\theta} J(\theta)$ is the gradient of cost function $J(\theta)$.

derivation of gradient

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(X^{(i)}) - y^{(i)})^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \sum_{i=1}^m (h_{\theta}(X^{(i)}) - y^{(i)}) \frac{\partial}{\partial \theta_j} (\theta^T X^{(i)} - y^{(i)})$$

Gradient

$$\frac{\partial}{\partial \theta_j} J(\theta) = \sum_{i=1}^m (h_{\theta}(X^{(i)}) - y^{(i)}) X_j^{(i)}$$

$$\nabla_{\theta} J(\theta) = \left[\frac{\partial}{\partial \theta_0} J(\theta) \quad \frac{\partial}{\partial \theta_1} J(\theta) \quad \cdots \quad \frac{\partial}{\partial \theta_n} J(\theta) \right]^T$$

stochastic gradient method

Data: α , m training examples $\{X^{(i)}, y^{(i)}\}_{i=1}^m$

Result: θ parameters

$\theta := \mathbf{0}$;

for $i = 1 \cdots m$ **do**

for $j = 0 \cdots n$ **do**

$\theta := \theta - \alpha(h_{\theta}(X^{(i)}) - y^{(i)})X_j^{(i)}$;

end

end

Algorithm 1: Stochastic gradient algorithm

closed form: add derivation involving trace

closed form solution

Solution

$$\theta = (X^T X)^{-1} X^T y$$

where X is an $m \times (n + 1)$ data matrix with $X[:, 1] = 1$ and y is a $m \times 1$ vector of labels

logistic regression

What's the problem?

- Given a set of m training examples $\{X^{(i)}, y^{(i)}\}_{i=1}^m$ where $X^{(i)} \in \mathbb{R}^{n+1}$ with $X_0^{(i)} = 1$ and $y^{(i)} \in \{0, 1\}$ is a label.
- Learn a hypothesis function $h(\cdot)$ such that for any new test example $x^{(t)}$, label for it can be predicted as $y^{(t)} = h(x^{(t)})$.
- In logistic regression, we assume:

$$h(x) = \text{sigmoid}(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

property of sigmoid function

$$\text{Let } g(x) = \frac{1}{1 + e^{-x}}$$

$$\begin{aligned}\text{Then, } g'(x) &= -\frac{1}{(1 + e^{-x})^2} \frac{d}{dx}(1 + e^{-x}) \\ &= \frac{e^{-x}}{(1 + e^{-x})^2} \\ &= \left(\frac{1}{1 + e^{-x}} \right) \left(\frac{e^{-x}}{1 + e^{-x}} \right) \\ &= \left(\frac{1}{1 + e^{-x}} \right) \left(1 - \frac{1}{1 + e^{-x}} \right)\end{aligned}$$

Thus,

$$g'(x) = g(x)(1 - g(x)) \quad (1)$$

optimization—maximize likelihood of the parameters θ

Assume

$$p(y = 1|x; \theta) = h_{\theta}(x)$$

$$p(y = 0|x; \theta) = 1 - h_{\theta}(x)$$

or more generally

$$p(y|x; \theta) = (h_{\theta}(x))^y (1 - h_{\theta}(x))^{(1-y)}$$

Likelihood of parameters

$$\begin{aligned} L(\theta) &= p(\mathbf{y}|\mathbf{X}; \theta) = \prod_{i=1}^m p(y^{(i)}|X^{(i)}; \theta) \\ &= \prod_{i=1}^m (h_{\theta}(X^{(i)}))^{y^{(i)}} (1 - h_{\theta}(X^{(i)}))^{(1-y^{(i)})} \end{aligned}$$

or equivalently maximize loglikelihood of the parameters θ

It's easier to maximize loglikelihood $\log(L(\theta))$ which is the same as maximizing likelihood $L(\theta)$

Optimization

$$\begin{aligned} l(\theta) &= \log L(\theta) \\ &= \sum_{i=1}^m [y^{(i)} \log h_{\theta}(X^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(X^{(i)}))] \\ \max_{\theta} l(\theta) \end{aligned}$$

Gradient ascent

Update rule:

$$\theta := \theta + \alpha \nabla_{\theta} l(\theta)$$

derivation of gradient of $l(\theta)$

$$\begin{aligned}\frac{\partial}{\partial \theta_j} l(\theta) &= \sum_{i=1}^m [y^{(i)} \log h_{\theta}(X^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(X^{(i)}))] \\&= \sum_{i=1}^m [y^{(i)} \frac{1}{h_{\theta}(X^{(i)})} \frac{\partial}{\partial \theta_j} h_{\theta}(X^{(i)}) \\&\quad + (1 - y^{(i)}) \frac{1}{(1 - h_{\theta}(X^{(i)}))} \frac{\partial}{\partial \theta_j} (1 - h_{\theta}(X^{(i)}))] \\&= \sum_{i=1}^m [y^{(i)} \frac{1}{h_{\theta}(X^{(i)})} h'_{\theta}(X^{(i)}) \frac{\partial}{\partial \theta_j} \theta^T X^{(i)} \\&\quad + (1 - y^{(i)}) \frac{1}{(1 - h_{\theta}(X^{(i)}))} (-h'_{\theta}(X^{(i)})) \frac{\partial}{\partial \theta_j} \theta^T X^{(i)}]\end{aligned}$$

derivation of gradient of $l(\theta)$

Using Eq. 1,

$$\begin{aligned}\frac{\partial}{\partial \theta_j} l(\theta) &= \sum_{i=1}^m \left[y^{(i)} \frac{1}{h_{\theta}(X^{(i)})} h_{\theta}(X^{(i)}) (1 - h_{\theta}(X^{(i)})) \frac{\partial}{\partial \theta_j} \theta^T X^{(i)} \right. \\ &\quad \left. + (1 - y^{(i)}) \frac{1}{(1 - h_{\theta}(X^{(i)}))} (-h_{\theta}(X^{(i)}) (1 - h_{\theta}(X^{(i)}))) \frac{\partial}{\partial \theta_j} \theta^T X^{(i)} \right] \\ &= \sum_{i=1}^m [y^{(i)} (1 - h_{\theta}(X^{(i)})) X_j^{(i)} + (1 - y^{(i)}) (-h_{\theta}(X^{(i)})) X_j^{(i)}]\end{aligned}$$

Derivative

$$\frac{\partial}{\partial \theta_j} l(\theta) = \sum_{i=1}^m (y^{(i)} - h_{\theta}(X^{(i)})) X_j^{(i)}$$